
LEADREACH AI IFSC CORP.

Source Verification & Token Budget Ledger

AI/LLM Token Usage & Cost Tracking

ENTITY	AI PROVIDER
LeadReach AI IFSC Corp., Ontario Corporation	Zhipu AI GLM

PRIMARY MODEL	FALLBACK MODEL
GLM-4.6v-Flash	GLM-4.7-Flash

REPORTING PERIOD	DOCUMENT VERSION
Monthly (Current)	1.4

Monthly Token Budget	Current Utilization	Monthly LLM Spend	Annual LLM Budget

CONFIDENTIAL

INVESTOR-FACING DOCUMENT DO NOT DISTRIBUTE

1. Executive Summary

Metric	Detail
Total Monthly Token Budget	45,000,000 tokens (across all agents)
Current Monthly Utilization	31,500,000 tokens (70% utilization)
Cost Per 1,000 Tokens (GLM-4.6v-flash)	\$0.0015
Cost Per 1,000 Tokens (GLM-4.7-flash)	\$0.0010
Monthly LLM Spend (Blended)	\$47,250
Annual LLM Budget	\$567,000

2. Token Budget Allocation by Agent (Monthly)

Agent	Budget (tokens)	Actual (tokens)	Utilization	Cost
Orchestrator	2,000,000	1,400,000	70%	\$2,100
Prospect Discovery	12,000,000	9,600,000	80%	\$14,400
Data Enrichment	8,000,000	5,600,000	70%	\$8,400
Web Research	8,000,000	6,400,000	80%	\$9,600
Lead Qualification	5,000,000	3,500,000	70%	\$5,250
Outreach Composer	5,000,000	3,000,000	60%	\$4,500
Pipeline Manager	2,000,000	1,000,000	50%	\$1,500
Report Generator	3,000,000	1,000,000	33%	\$1,500
TOTAL	45,000,000	31,500,000	70%	\$47,250

3. Token Usage by Skill

Skill	Agent	Avg Tokens/Call	Monthly Calls	Monthly Tokens	Cost
discover_companies	Prospect Discovery	2,500	2,400	6,000,000	\$9,000
deep_search	Prospect Discovery	3,000	1,200	3,600,000	\$5,400
enrich_leads	Data Enrichment	1,800	2,000	3,600,000	\$5,400
extract_contacts	Data Enrichment	1,200	1,667	2,000,000	\$3,000
research_company	Web Research	2,200	1,800	3,960,000	\$5,940
market_analysis	Web Research	2,800	871	2,440,000	\$3,660
score_bant	Lead Qualification	800	1,500	1,200,000	\$1,800
score_meddico	Lead Qualification	1,200	1,000	1,200,000	\$1,800
score_icp	Lead Qualification	1,000	1,100	1,100,000	\$1,650
compose_email	Outreach Composer	1,500	1,200	1,800,000	\$2,700
compose_linkedin	Outreach Composer	1,200	1,000	1,200,000	\$1,800
design_sequence	Outreach Composer	2,000	500	1,000,000	\$1,500

4. Model Routing Distribution

Model	Primary Calls	Fallback Calls	Total Tokens	% of Budget	Cost
glm-4.6v-flash	85%	-	26,775,000	59.5%	\$40,163
glm-4.7-flash	15% (fallback)	100%	4,725,000	10.5%	\$4,725
TOTAL	100%	-	31,500,000	70%	\$44,888

5. Thinking Budget Allocation

Budget Level	Token Limit	Use Cases	Allocation
Quick	512	Intent classification, health checks, routing	10%
Standard	2,048	Chat, data extraction, scoring	65%
Deep	4,096	Deep research, multi-step analysis, complex reasoning	25%

6. Rate Limiting & Resilience

Global rate limit: 1.5 s minimum interval + 500 ms jitter between calls

Model health tracking: 429 cooldown - 1 minute after 2 consecutive 429 responses

Automatic backoff: 3 s x attempt for rate limits; 2 s x attempt for gateway errors

3-tier fallback: Real search -> LLM knowledge generation -> Hardcoded industry knowledge

7. Cost Optimization Strategies

Multi-model routing reduces cost by 15% vs. single-model deployment

Thinking budget allocation prevents unnecessary deep reasoning on simple tasks

Rate limiting and backoff prevent API waste on transient errors

Agent memory system reduces redundant API calls by 20-30% over time

Budgeted at 70% utilization to allow for growth without exceeding limits

8. 12-Month Token Budget Projection

Period	Projected Users	Token Budget	Projected Spend
Months 1-3	86	45,000,000	\$47,250/mo
Months 4-6	172	60,000,000	\$63,000/mo
Months 7-9	280	85,000,000	\$89,250/mo
Months 10-12	345	105,000,000	\$110,250/mo
Annual Total	-	-	\$927,000