

LEADREACH AI CORP.

Source Verification & Token Budget Ledger

AI/LLM Token Usage & Cost Tracking

ENTITY

LeadReach AI Corp., Delaware C-Corp

AI PROVIDER

Zhipu AI GLM

PRIMARY MODEL

GLM-4.6v-Flash

FALLBACK MODEL

GLM-4.7-Flash

REPORTING PERIOD

Monthly (Current)

DOCUMENT VERSION

1.4

45M

Monthly Token Budget

70%

Current Utilization

\$47.3K

Monthly LLM Spend

\$567K

Annual LLM Budget

CONFIDENTIAL

1. Executive Summary

45M Monthly Token Budget	70% Current Utilization	\$47,250 Monthly LLM Spend	\$567K Annual LLM Budget
------------------------------------	-----------------------------------	--------------------------------------	------------------------------------

Metric	Detail
Total Monthly Token Budget	45,000,000 tokens (across all agents)
Current Monthly Utilization	31,500,000 tokens (70% utilization)
Cost Per 1,000 Tokens (GLM-4.6v-flash)	\$0.0015
Cost Per 1,000 Tokens (GLM-4.7-flash)	\$0.0010
Monthly LLM Spend (Blended)	\$47,250
Annual LLM Budget	\$567,000

2. Token Budget Allocation by Agent (Monthly)

Agent	Budget (tokens)	Actual (tokens)	Utilization	Cost
Orchestrator	2,000,000	1,400,000	70%	\$2,100
Prospect Discovery	12,000,000	9,600,000	80%	\$14,400
Data Enrichment	8,000,000	5,600,000	70%	\$8,400
Web Research	8,000,000	6,400,000	80%	\$9,600
Lead Qualification	5,000,000	3,500,000	70%	\$5,250
Outreach Composer	5,000,000	3,000,000	60%	\$4,500
Pipeline Manager	2,000,000	1,000,000	50%	\$1,500
Report Generator	3,000,000	1,000,000	33%	\$1,500
TOTAL	45,000,000	31,500,000	70%	\$47,250

3. Token Usage by Skill

Skill	Agent	Avg Tokens/Call	Monthly Calls	Monthly Tokens	Cost
discover_companies	Prospect Discovery	2,500	2,400	6,000,000	\$9,000
deep_search	Prospect Discovery	3,000	1,200	3,600,000	\$5,400
enrich_leads	Data Enrichment	1,800	2,000	3,600,000	\$5,400
extract_contacts	Data Enrichment	1,200	1,667	2,000,000	\$3,000
research_company	Web Research	2,200	1,800	3,960,000	\$5,940
market_analysis	Web Research	2,800	871	2,440,000	\$3,660
score_bant	Lead Qualification	800	1,500	1,200,000	\$1,800
score_meddic	Lead Qualification	1,200	1,000	1,200,000	\$1,800

score_icp	Lead Qualification	1,000	1,100	1,100,000	\$1,650
compose_email	Outreach Composer	1,500	1,200	1,800,000	\$2,700
compose_linkedin	Outreach Composer	1,200	1,000	1,200,000	\$1,800
design_sequence	Outreach Composer	2,000	500	1,000,000	\$1,500

4. Model Routing Distribution

Model	Primary Calls	Fallback Calls	Total Tokens	% of Budget	Cost
glm-4.6v-flash	85%	—	26,775,000	59.5%	\$40,163
glm-4.7-flash	15% (fallback)	100%	4,725,000	10.5%	\$4,725
TOTAL	100%	—	31,500,000	70%	\$44,888

5. Thinking Budget Allocation

Budget Level	Token Limit	Use Cases	Allocation
Quick	512	Intent classification, health checks, routing	10%
Standard	2,048	Chat, data extraction, scoring	65%
Deep	4,096	Deep research, multi-step analysis, complex reasoning	25%

6. Rate Limiting & Resilience

- Global rate limit: 1.5 s minimum interval + 500 ms jitter between calls
- Model health tracking: 429 cooldown — 1 minute after 2 consecutive 429 responses
- Automatic backoff: 3 s × attempt for rate limits; 2 s × attempt for gateway errors
- 3-tier fallback: Real search → LLM knowledge generation → Hardcoded industry knowledge

7. Cost Optimization Strategies

- Multi-model routing reduces cost by 15% vs. single-model deployment
- Thinking budget allocation prevents unnecessary deep reasoning on simple tasks
- Rate limiting and backoff prevent API waste on transient errors
- Agent memory system reduces redundant API calls by 20–30% over time
- Budgeted at 70% utilization to allow for growth without exceeding limits

8. 12-Month Token Budget Projection

Period	Projected Users	Token Budget	Projected Spend
Months 1–3	86	45,000,000	\$47,250/mo

Months 4-6	172	60,000,000	\$63,000/mo
Months 7-9	280	85,000,000	\$89,250/mo
Months 10-12	345	105,000,000	\$110,250/mo
Annual Total	—	—	\$927,000